

# CLB-SVD V2.0

March 18, 2024

```
[1]: import pandas as pd
import numpy as np
from scipy.sparse import csc_matrix
from scipy.sparse.linalg import svds
from sklearn.metrics import mean_squared_error
from math import sqrt

[2]: #Pandas work with data in a table format [for data manipulation and analysis]
#NumPy is support for arrays and mathematical functions. [numerical computing]
# scipy.sparse.linalg.svds is a function to perform Singular Value
    ↳Decomposition (SVD) a way to decompose a matrix.
#sklearn.metrics.mean_squared_error calculates the mean squared error, a
    ↳measure of how close predictions are to the actual outcomes.
#sqrt from the math module calculates the square root.

[3]: # Parsing the data
def parse_data(file_path):
    data = []
    current_movie_id = None
    with open(file_path, 'r', encoding='utf-8-sig') as file:
        for line in file:
            try:
                if ':' in line:
                    current_movie_id = int(line.split(':')[0])
                else:
                    customer_id, rating, _ = line.strip().split(',')
                    data.append([current_movie_id, int(customer_id),
    ↳int(rating)])
            except ValueError:
                # This handles lines that don't have the expected format
                continue # Skips to the next line
    return pd.DataFrame(data, columns=['MovieID', 'CustomerID', 'Rating'])

# Function to load data for training, test, and validation
def load_dataset(file_paths):
    data_frames = [parse_data(file_path) for file_path in file_paths]
    ↳#parse_data function to load and structure the data into a DataFrame
```

```

combined_data = pd.concat(data_frames)
return combined_data

# File paths setup
train_files = ['training_set_c1.txt', 'training_set_c2.txt']
test_file = ['test_set_c1.txt', 'test_set_c2.txt']
validation_file = ['validation_set_c1.txt', 'validation_set_c2.txt']

# Load datasets
train_movie_data = load_dataset(train_files)

train_movie_data.head()

```

```

[3]:
  MovieID  CustomerID  Rating
0        1        401047      4
1        1        14756      4
2        1       2566259      5
3        1       1398626      2
4        1       1294335      2

```

```

[4]: # Converts the DataFrame into a user-item matrix where rows represent
      ↪ customers, columns represent movies, and values are ratings.
      # Missing ratings are filled with 0, indicating unrated movies.

```

```

[5]: # Create Training user-item matrix
train_ratings_df = train_movie_data.pivot(index='CustomerID',
      ↪ columns='MovieID', values='Rating').fillna(0)

# Convert the DataFrame to a Compressed Sparse Column (CSC) matrix
train_ratings_matrix = csc_matrix(train_ratings_df.values)

num_users, num_movies = train_ratings_matrix.shape
print(f"Number of users: {num_users}, Number of movies: {num_movies}")

```

/tmp/ipykernel\_15519/1906747628.py:2: PerformanceWarning: The following operation may generate 4384890210 cells in the resulting pandas object.

```

train_ratings_df = train_movie_data.pivot(index='CustomerID',
columns='MovieID', values='Rating').fillna(0)

```

Number of users: 476101, Number of movies: 9210

```

[6]: #Performs SVD on the sparse matrix, decomposing it into matrices U, Sigma, and
      ↪ V^T. Sigma is converted into a diagonal matrix.

```

```

[7]: # "k" value represents the number of singular values (or latent factors)
      k=11
      # Perform SVD with the updated k value

```

```
U, sigma, Vt = svds(train_ratings_matrix, k=k)
sigma = np.diag(sigma)
```

[8]: *## SVD is a technique used to break down the ratings matrix into three smaller matrices.*  
*#It helps to identify patterns in the data, such as similar preferences among customers.*  
*## k is the number of latent factors you want to keep. These factors are like hidden themes that capture the essence of the data.*

[9]: **def** predict(matrix, U, sigma, Vt):  
    *# Calculate mean user rating with correct reshaping*  
    mean\_user\_rating = matrix.mean(axis=1).reshape(-1, 1)  
    *# Ensure that U, sigma, and Vt are correctly aligned with the input matrix dimensions*  
    preds = np.dot(np.dot(U, sigma), Vt) + mean\_user\_rating  
    **return** preds

[10]: *#This line reconstructs the ratings matrix using the decomposed matrices obtained from SVD, adjusted by the average rating for each user.*  
*#It essentially predicts the ratings that users might give to movies they haven't rated.*

[11]: *# RMSE calculation*  
**def** calculate\_rmse(actual, predicted):  
    *# Ensure both actual and predicted are numpy arrays*  
    mask = actual.nonzero()  
    actual\_filtered = actual[mask].flatten()  
    predicted\_filtered = predicted[mask].flatten()  
    **return** sqrt(mean\_squared\_error(actual\_filtered, predicted\_filtered))  
  
*# Assuming train\_ratings\_matrix is your actual ratings and train\_preds is your predictions*  
train\_preds = predict(train\_ratings\_df.values, U, sigma, Vt)  
*# Calculate RMSE*  
**print**('Training RMSE:', calculate\_rmse(train\_ratings\_matrix.toarray(), train\_preds))

Training RMSE: 2.8289205193008806

[12]: *#RMSE (Root Mean Square Error) is a way to measure how accurate the predictions are. It calculates the difference between the predicted ratings and the actual ratings given by the users, providing a single number that represents the average error.*  
*#A lower RMSE means the predictions are closer to the actual ratings, indicating a better performing model.*

```
[13]: test_movie_data = parse_data(test_file[0])
```

```
[14]: # Create a Test user-item matrix
test_ratings_df = test_movie_data.pivot(index='CustomerID', columns='MovieID',
    ↪ values='Rating').reindex(index=train_ratings_df.index,
    ↪ columns=train_ratings_df.columns).fillna(0)

# Convert to CSC format
test_ratings_matrix = csc_matrix(test_ratings_df.values)

num_users, num_movies = test_ratings_matrix.shape
print(f"Number of users: {num_users}, Number of movies: {num_movies}")
```

Number of users: 476101, Number of movies: 9210

```
[15]: # Make predictions with adjusted dimensions
test_preds = predict(test_ratings_df.values, U, sigma, Vt)
```

```
[16]: print("Shape of test_matrix.values:", test_ratings_df.values.shape)
print("Shape of test_preds:", test_preds.shape)
```

Shape of test\_matrix.values: (476101, 9210)

Shape of test\_preds: (476101, 9210)

```
[17]: def calculate_rmse(actual, predicted):
    mask = actual > 0 # Assuming ratings are positive and 0 indicates missing
    ↪ rating
    print(f"Non-zero entries in 'actual': {np.sum(mask)}")

    actual_filtered = actual[mask]
    predicted_filtered = predicted[mask]
    print(f"Shape of 'actual': {actual.shape}")
    print(f"Shape of 'predicted': {predicted.shape}")

    return sqrt(mean_squared_error(actual_filtered, predicted_filtered))

# Calculate RMSE
print('Test RMSE:', calculate_rmse(test_ratings_matrix.toarray(), test_preds))
```

Non-zero entries in 'actual': 3607424

Shape of 'actual': (476101, 9210)

Shape of 'predicted': (476101, 9210)

Test RMSE: 2.948704925212078

```
[18]: validation_movie_data = parse_data(validation_file[0])
```

```
[19]: # Create a validation matrix similar to the training matrix
validation_ratings_df = validation_movie_data.pivot(index='CustomerID',
↳columns='MovieID', values='Rating').reindex(index=train_ratings_df.index,
↳columns=train_ratings_df.columns).fillna(0)
# Convert to CSC format
validation_ratings_matrix = csc_matrix(validation_ratings_df.values)

num_users, num_movies = validation_ratings_matrix.shape
print(f"Number of users: {num_users}, Number of movies: {num_movies}")
```

Number of users: 476101, Number of movies: 9210

```
[20]: # Make predictions with adjusted dimensions
validation_preds = predict(validation_ratings_df.values, U, sigma, Vt)
```

```
[21]: print("Shape of validation_matrix.values:", validation_ratings_df.values.shape)
print("Shape of validation_preds:", validation_preds.shape)
```

Shape of validation\_matrix.values: (476101, 9210)

Shape of validation\_preds: (476101, 9210)

```
[22]: print('validation RMSE:', calculate_rmse(validation_ratings_matrix.toarray(),
↳validation_preds))
```

Non-zero entries in 'actual': 3607427

Shape of 'actual': (476101, 9210)

Shape of 'predicted': (476101, 9210)

validation RMSE: 2.9491639886583707

[ ]:

[ ]:

[ ]: