

# SHERPA\_V2.0

March 18, 2024

```
[1]: import pandas as pd
import numpy as np
from scipy.sparse import csc_matrix
from numpy.linalg import lstsq
from sklearn.metrics import mean_squared_error
from math import sqrt
```

## 0.1 Dataset Importing Function

First, let's include the function for importing the dataset, which parses the data from a given file and formats it into a pandas DataFrame.

```
[2]: movies_df = pd.read_csv("movie_titles.csv", on_bad_lines='skip')
```

```
[3]: # Parsing the data
def parse_data(file_path):
    data = []
    current_movie_id = None
    with open(file_path, 'r', encoding='utf-8-sig') as file:
        for line in file:
            try:
                if ':' in line:
                    current_movie_id = int(line.split(':')[0])
                else:
                    customer_id, rating, _ = line.strip().split(',')
                    data.append([current_movie_id, int(customer_id),
↪int(rating)])
            except ValueError:
                # This handles lines that don't have the expected format
                continue # Skips to the next line
    return pd.DataFrame(data, columns=['MovieID', 'CustomerID', 'Rating'])

# Function to load data for training, test, and validation
def load_dataset(file_paths):
    data_frames = [parse_data(file_path) for file_path in file_paths]
    ↪#parse_data function to load and structure the data into a DataFrame
    combined_data = pd.concat(data_frames)
```

```

    return combined_data

# File paths setup
train_files = ['training_set_c1.txt', 'training_set_c2.txt']
test_file = ['test_set_c1.txt', 'test_set_c2.txt']
validation_file = ['validation_set_c1.txt', 'validation_set_c2.txt']

# Load datasets
train_movie_data = load_dataset(train_files)

train_movie_data.head()

```

```

[3]:   MovieID  CustomerID  Rating
      0         1      401047      4
      1         1      14756      4
      2         1     2566259      5
      3         1     1398626      2
      4         1     1294335      2

```

## 0.2 Preparing the Ratings Matrix

After importing the dataset, convert it to a sparse matrix format that the ALS algorithm can process.

```

[4]: from scipy.sparse import csr_matrix

# Create a user-item matrix
train_ratings_df = train_movie_data.pivot(index='CustomerID',
      ↪columns='MovieID', values='Rating').fillna(0)

# Convert to CSR format
train_ratings_matrix = csr_matrix(train_ratings_df.values)

```

/tmp/ipykernel\_12920/1423468332.py:4: PerformanceWarning: The following operation may generate 4384890210 cells in the resulting pandas object.

```

train_ratings_df = train_movie_data.pivot(index='CustomerID',
columns='MovieID', values='Rating').fillna(0)

```

## 0.3 Optimized ALS with Sparse Matrices

```

[5]: import numpy as np
      from scipy.sparse import csr_matrix
      from numpy.linalg import lstsq

      def update_U(M, U, lambda_reg, ratings):
          """
          Update user features matrix U.

```

```

:param M: Movie features matrix.
:param U: User features matrix to update.
:param lambda_reg: Regularization parameter.
:param ratings: Ratings CSR matrix.
:return: Updated user features matrix.
"""
num_factors = U.shape[1] # Number of latent factors
lambda_I = lambda_reg * np.eye(num_factors)
updated_U = np.zeros(U.shape)

for i in range(U.shape[0]):
    userRatedIndices = ratings[i].nonzero()[1]
    M_sub = M[userRatedIndices, :]
    ratings_sub = ratings[i, userRatedIndices].toarray().flatten()

    # Constructing the system to solve
    A = M_sub.T @ M_sub + lambda_I * len(userRatedIndices)
    b = M_sub.T @ ratings_sub
    updated_U[i, :] = lstsq(A, b, rcond=None)[0]

return updated_U

def update_M(M, U, lambda_reg, ratings):
    """
    Update movie features matrix M.

    :param M: Movie features matrix to update.
    :param U: User features matrix.
    :param lambda_reg: Regularization parameter.
    :param ratings: Ratings CSR matrix.
    :return: Updated movie features matrix.
    """
    num_factors = M.shape[1] # Number of latent factors
    lambda_I = lambda_reg * np.eye(num_factors)
    updated_M = np.zeros(M.shape)

    for i in range(M.shape[0]):
        movieRatedByIndices = ratings[:, i].nonzero()[0]
        U_sub = U[movieRatedByIndices, :]
        ratings_sub = ratings[movieRatedByIndices, i].toarray().flatten()

        # Constructing the system to solve
        A = U_sub.T @ U_sub + lambda_I * len(movieRatedByIndices)
        b = U_sub.T @ ratings_sub
        updated_M[i, :] = lstsq(A, b, rcond=None)[0]

```

```

    return updated_M

def ALS(ratings, num_factors=50, lambda_reg=0.1, iterations=10):
    """
    Alternating Least Squares algorithm for matrix factorization.

    :param ratings: Original ratings matrix as CSR.
    :param num_factors: Number of latent factors.
    :param lambda_reg: Regularization parameter.
    :param iterations: Number of iterations to run ALS.
    :return: User and Movie latent features matrices.
    """
    num_users, num_movies = ratings.shape

    # Initialize user and movie matrices with small random values
    U = np.random.rand(num_users, num_factors) * 0.01
    M = np.random.rand(num_movies, num_factors) * 0.01

    for iteration in range(iterations):
        U = update_U(M, U, lambda_reg, ratings)
        M = update_M(M, U, lambda_reg, ratings)

    return U, M

```

## 0.4 Running the ALS Algorithm

Now, you can use the previously defined ALS algorithm, ensuring it's ready to process the CSR matrix created from your dataset.

```

[6]: # Assuming the ALS function and update functions are defined as in the previous
    ↪ example

```

```

num_factors = 11 # Adjust based on your preference
lambda_reg = 0.1 # Regularization parameter
iterations = 5 # Number of iterations for ALS

# Running ALS on the ratings matrix
U, M = ALS(train_ratings_matrix, num_factors=num_factors,
    ↪ lambda_reg=lambda_reg, iterations=iterations)

# U and M are the user and item (movie) latent factors matrices

```

```

[21]: # Generate predictions
train_predictions = U.dot(M.T)

train_preds_df = pd.DataFrame(train_predictions, columns=train_ratings_df.
    ↪ columns, index=train_ratings_df.index)

```

```
[22]: #Content-Based Filtering
```

```
[23]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import sigmoid_kernel
from sklearn.metrics import ndcg_score

movies_df = movies_df.fillna('')
# Function to create weighted text
def create_weighted_text(row):
    return (row['Overview'] + ' ') * 45 + (row['Genre'] + ' ') * 25 + \
        (row['Director'] + ' ') * 15 + (row['Cast'] + ' ') * 15
movies_df['weighted_text'] = movies_df.apply(create_weighted_text, axis=1)

# Initialize TF-IDF Vectorizer
# from sklearn.feature_extraction.text import TfidfVectorizer
tfv = TfidfVectorizer(min_df=3, max_features=None, strip_accents='unicode',
                      analyzer='word', token_pattern=r'\w{1,}',
                      ngram_range=(1,3), stop_words='english')
# Fit the TF-IDF on the 'weighted_text'
tfv_matrix = tfv.fit_transform(movies_df['weighted_text'])
sig = sigmoid_kernel(tfv_matrix, tfv_matrix)
# Reverse mapping of indices and movie titles
indices = pd.Series(movies_df.index, index=movies_df['Movie_Name']).
    ↪drop_duplicates()
```

```
[24]: #Hybrid Recommendation System
```

```
[41]: def hybrid_recommendations(user_id=None, movie_name=None, preds_df=None,
    ↪movies_df=movies_df, sig=sig, indices=indices, top_n=10):
    if preds_df is None:
        raise ValueError("The predictions dataframe (preds_df) is required.")

    # Initialize recommendation lists
    final_recs = []

    # Fetch Content-Based Recommendations
    content_based_recs = []
    if movie_name in indices:
        idx = indices[movie_name]
        sig_scores = list(enumerate(sig[idx]))
        sig_scores = sorted(sig_scores, key=lambda x: x[1], reverse=True)
        movie_indices = [i[0] for i in sig_scores[1:top_n+1]]
        content_based_recs = movies_df.iloc[movie_indices]['Movie_Name'].
            ↪tolist()

    # For existing users with a search query, combine collaborative and
    ↪content-based recommendations
```

```

    if user_id and movie_name:
        # Fetch collaborative filtering recommendations based on historical
        ↪ ratings
        collaborative_recs_ids = preds_df.loc[user_id].
        ↪ sort_values(ascending=False).head(top_n * 2).index.tolist()
        collaborative_recs_names = movies_df[movies_df['MovieID'].
        ↪ isin(collaborative_recs_ids)]['Movie_Name'].tolist()

        # Combine lists with simple deduplication, prioritizing content-based
        ↪ recommendations
        seen = set(content_based_recs)
        combined_recs = content_based_recs + [rec for rec in
        ↪ collaborative_recs_names if rec not in seen]

        # Limit to top_n recommendations after combining
        final_recs = combined_recs[:top_n]
    elif user_id:
        # Only collaborative recommendations for existing users without search
        ↪ query
        collaborative_recs_ids = preds_df.loc[user_id].
        ↪ sort_values(ascending=False).head(top_n).index.tolist()
        final_recs = movies_df[movies_df['MovieID'].
        ↪ isin(collaborative_recs_ids)]['Movie_Name'].tolist()
    else:
        # Only content-based recommendations for new users with a search query
        final_recs = content_based_recs

    return final_recs

# Testing the function with your scenarios
user_id = 401047 # Example user ID
movie_name = "The Company" # Example movie name

print("Collaborative Recommendations for Existing User (No Search):")
collab_recs = hybrid_recommendations(user_id=user_id, preds_df=train_preds_df,
    ↪ top_n=10)
for movie in collab_recs:
    print(movie)

print("\nHybrid Recommendations for Existing User (With Search):")
hybrid_recs = hybrid_recommendations(user_id=user_id, movie_name=movie_name,
    ↪ preds_df=train_preds_df, top_n=10)
for movie in hybrid_recs:
    print(movie)

print("\nContent-Based Recommendations for New User (With Search):")

```

```

content_recs = hybrid_recommendations(movie_name=movie_name,
    ↪preds_df=train_preds_df, top_n=10)
for movie in content_recs:
    print(movie)

```

Collaborative Recommendations for Existing User (No Search):

ABC Primetime: Mel Gibson's The Passion of the Christ  
 Armageddon  
 Coach Carter  
 Secondhand Lions  
 The Winds of War  
 Braveheart  
 In the Face of Evil: Reagan's War in Word and Deed  
 Pretty Woman  
 24: Season 1  
 CSI: Season 3

Hybrid Recommendations for Existing User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City  
 A Raisin in the Sun  
 Robin and the 7 Hoods  
 Swan Lake: Tchaikovsky (Matthew Bourne)  
 Out of Sync  
 Orchestra Rehearsal  
 Category 6: Day of Destruction  
 What Have I Done to Deserve This?

Content-Based Recommendations for New User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City  
 A Raisin in the Sun  
 Robin and the 7 Hoods  
 Swan Lake: Tchaikovsky (Matthew Bourne)  
 Out of Sync  
 Orchestra Rehearsal  
 Category 6: Day of Destruction  
 What Have I Done to Deserve This?

```

[39]: from sklearn.metrics import mean_squared_error
      from math import sqrt

      def calculate_rmse(actual, predictions):
          mask = actual.nonzero() # Only consider non-zero entries
          actual = actual[mask]

```

```

    predictions = predictions[mask]
    return sqrt(mean_squared_error(actual, predictions))

# Convert the predictions matrix to a dense format since ratings_matrix is
↳ sparse
train_ratings_matrix = train_ratings_matrix.toarray()

# Calculate RMSE
rmse = calculate_rmse(train_ratings_matrix.toarray(), train_predictions)
print('Training RMSE:', rmse)

```

Training RMSE: 0.8771686076210009

```

[40]: # Load datasets
test_movie_data = load_dataset(test_file)

# Create a Train user-item matrix
test_ratings_df= test_movie_data.pivot(index='CustomerID', columns='MovieID',
↳ values='Rating').reindex(index=train_ratings_df.index,
↳ columns=train_ratings_df.columns).fillna(0)

# Convert to CSR format
test_ratings_matrix = csr_matrix(test_ratings_df.values)

```

/tmp/ipykernel\_12920/2213329239.py:5: PerformanceWarning: The following operation may generate 4087377814 cells in the resulting pandas object.

```

test_ratings_df= test_movie_data.pivot(index='CustomerID', columns='MovieID',
values='Rating').reindex(index=train_ratings_df.index,
columns=train_ratings_df.columns).fillna(0)

```

```

[43]: # Mapping test user and movie indices to training set indices
test_user_indices = [np.where(train_ratings_df.index == uid)[0][0] for uid in
↳ test_ratings_df.index if uid in train_ratings_df.index]
test_movie_indices = [np.where(train_ratings_df.columns == mid)[0][0] for mid
↳ in test_ratings_df.columns if mid in train_ratings_df.columns]

# Generate predictions for test set
test_predictions = U[test_user_indices, :] @ M.T[:, test_movie_indices]

# Generate predictions
test_preds_df = pd.DataFrame(test_predictions, columns=test_ratings_df.columns,
↳ index=test_ratings_df.index)

```

```

[55]: # Testing the function with your scenarios
user_id = 401047 # Example user ID
movie_name = "The Company" # Example movie name

```



```

print("Collaborative Recommendations for Existing User (No Search):")
collab_recs = hybrid_recommendations(user_id=user_id, preds_df=test_preds_df,
    ↪top_n=10)
for movie in collab_recs:
    print(movie)

print("\nHybrid Recommendations for Existing User (With Search):")
hybrid_recs = hybrid_recommendations(user_id=user_id, movie_name=movie_name,
    ↪preds_df=test_preds_df, top_n=10)
for movie in hybrid_recs:
    print(movie)

print("\nContent-Based Recommendations for New User (With Search):")
content_recs = hybrid_recommendations(movie_name=movie_name,
    ↪preds_df=test_preds_df, top_n=10)
for movie in content_recs:
    print(movie)

```

Collaborative Recommendations for Existing User (No Search):

ABC Primetime: Mel Gibson's The Passion of the Christ  
 Armageddon  
 Coach Carter  
 Secondhand Lions  
 The Winds of War  
 Braveheart  
 In the Face of Evil: Reagan's War in Word and Deed  
 Pretty Woman  
 24: Season 1  
 CSI: Season 3

Hybrid Recommendations for Existing User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City  
 A Raisin in the Sun  
 Robin and the 7 Hoods  
 Swan Lake: Tchaikovsky (Matthew Bourne)  
 Out of Sync  
 Orchestra Rehearsal  
 Category 6: Day of Destruction  
 What Have I Done to Deserve This?

Content-Based Recommendations for New User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City  
 A Raisin in the Sun

Robin and the 7 Hoods  
Swan Lake: Tchaikovsky (Matthew Bourne)  
Out of Sync  
Orchestra Rehearsal  
Category 6: Day of Destruction  
What Have I Done to Deserve This?

```
[48]: # Calculate RMSE for test set
test_rmse = calculate_rmse(test_ratings_matrix.toarray(), test_predictions)
print('Test RMSE:', test_rmse)
```

Test RMSE: 0.9181483695344229

```
[49]: # Load datasets
validation_movie_data = load_dataset(validation_file)

# Create a user-item matrix
validation_ratings_df = validation_movie_data.pivot(index='CustomerID',
    ↳ columns='MovieID', values='Rating').reindex(index=train_ratings_df.index,
    ↳ columns=train_ratings_df.columns).fillna(0)

# Convert to CSR format
validation_ratings_matrix = csr_matrix(validation_ratings_df.values)
```

/tmp/ipykernel\_12920/1611650772.py:5: PerformanceWarning: The following operation may generate 4087812450 cells in the resulting pandas object.

```
validation_ratings_df = validation_movie_data.pivot(index='CustomerID',
columns='MovieID', values='Rating').reindex(index=train_ratings_df.index,
columns=train_ratings_df.columns).fillna(0)
```

```
[51]: # Mapping validation user and movie indices to training set indices
validation_user_indices = [np.where(train_ratings_df.index == uid)[0][0] for
    ↳ uid in validation_ratings_df.index if uid in train_ratings_df.index]
validation_movie_indices = [np.where(train_ratings_df.columns == mid)[0][0] for
    ↳ mid in validation_ratings_df.columns if mid in train_ratings_df.columns]

# Generate predictions for validation set
validation_predictions = U[validation_user_indices, :] @ M.T[:,
    ↳ validation_movie_indices]

validation_preds_df = pd.DataFrame(validation_predictions,
    ↳ columns=validation_ratings_df.columns, index=validation_ratings_df.index)
```

```
[56]: # Testing the function with your scenarios
user_id = 401047 # Example user ID
movie_name = "The Company" # Example movie name
```

```

print("Collaborative Recommendations for Existing User (No Search):")
collab_recs = hybrid_recommendations(user_id=user_id,
    ↪preds_df=validation_preds_df, top_n=10)
for movie in collab_recs:
    print(movie)

print("\nHybrid Recommendations for Existing User (With Search):")
hybrid_recs = hybrid_recommendations(user_id=user_id, movie_name=movie_name,
    ↪preds_df=validation_preds_df, top_n=10)
for movie in hybrid_recs:
    print(movie)

print("\nContent-Based Recommendations for New User (With Search):")
content_recs = hybrid_recommendations(movie_name=movie_name,
    ↪preds_df=validation_preds_df, top_n=10)
for movie in content_recs:
    print(movie)

```

Collaborative Recommendations for Existing User (No Search):

ABC Primetime: Mel Gibson's The Passion of the Christ  
 Armageddon  
 Coach Carter  
 Secondhand Lions  
 The Winds of War  
 Braveheart  
 In the Face of Evil: Reagan's War in Word and Deed  
 Pretty Woman  
 24: Season 1  
 CSI: Season 3

Hybrid Recommendations for Existing User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City  
 A Raisin in the Sun  
 Robin and the 7 Hoods  
 Swan Lake: Tchaikovsky (Matthew Bourne)  
 Out of Sync  
 Orchestra Rehearsal  
 Category 6: Day of Destruction  
 What Have I Done to Deserve This?

Content-Based Recommendations for New User (With Search):

Center Stage  
 Ballet Favorites  
 Expo: Magic of the White City

A Raisin in the Sun  
Robin and the 7 Hoods  
Swan Lake: Tchaikovsky (Matthew Bourne)  
Out of Sync  
Orchestra Rehearsal  
Category 6: Day of Destruction  
What Have I Done to Deserve This?

```
[53]: # Calculate RMSE for validation set
validation_rmse = calculate_rmse(validation_ratings_matrix.toarray(),
    ↪validation_predictions)
print('validation RMSE:', validation_rmse)
```

validation RMSE: 0.9183960844224338

```
[ ]: # 0.9525 Cinematch Netflix score
```

```
[ ]:
```

```
[ ]:
```